

NGUYEN KHAC BAO

Fullstack AI Integration Intern

Ho Chi Minh City, Vietnam | nguyenkhacbaowork564@gmail.com | (+84) 039 586 4429 | github.com/NguyenKhacBao564

CAREER OBJECTIVE

AI integration intern candidate focused on turning AI models into usable products through backend APIs, RAG pipelines, vector databases, Docker, and cloud deployment. Hands-on with FastAPI, Streamlit, Gemini API, Qdrant Cloud, Cloud SQL, and Cloud Run for deployable AI demos.

EDUCATION

Posts and Telecommunications Institute of Technology (PTIT)

2022 – 2027

Bachelor of Information Technology

Relevant Coursework: Machine Learning, Deep Learning, Computer Vision.

TECHNICAL SKILLS

AI Integration: RAG pipelines, Gemini API, embeddings, prompt engineering, OpenAI-compatible APIs, LLM endpoint integration

Backend / Cloud: FastAPI, REST API, WebSocket, Docker, Google Cloud Run, Cloud SQL, Secret Manager, AWS EC2, Vast.ai GPU

Databases: Qdrant Cloud, MySQL/Cloud SQL, SQL basics, JSON/CSV data processing

Frontend: Streamlit, browser webcam demo, basic HTML/JavaScript

ML / CV: PyTorch, YOLOv8n, OpenCV, OCR, multi-object tracking, dataset QA

Tools: Git/GitHub, Linux, Cloud Build, Artifact Registry

LANGUAGES

Vietnamese: Native | English: TOEIC 610/990; technical reading proficiency for documentation, model cards, API references, and engineering articles.

PROJECTS

Vietnamese Legal RAG Chatbot

2026

Personal Project / Python, FastAPI, Streamlit, Gemini API, Qdrant Cloud, Cloud SQL, Docker, Google Cloud Run / Source: [GitHub](#)

- Built a cloud-ready fullstack RAG product demo with a Streamlit frontend and FastAPI backend deployed as separate Google Cloud Run services.
- Integrated Gemini as the default LLM and embedding provider; indexed the full 32,980-record Vietnamese legal QA dataset into Qdrant Cloud for retrieval-based answers.
- Used Cloud SQL/MySQL for chat history and environment-based configuration with Secret Manager for API keys and database credentials.
- Designed optional extensions for async document indexing with Celery/Redis-compatible broker and fine-tuned Vietnamese legal LLM serving through an OpenAI-compatible endpoint.

Real-Time Face Mask Compliance Detection System

2026

Personal Project / Python, FastAPI, WebSocket, YOLOv8n, OpenCV, Docker, Google Cloud Run / Source: [GitHub](#)

- Built and deployed an AI inference service with FastAPI REST + WebSocket endpoints and a browser webcam frontend running over Cloud Run HTTPS/WSS.
- Containerized the service for Cloud Run and logged warm endpoint latency around 151 ms for REST and 127 ms for WSS in the demo environment.
- Trained a YOLOv8n 3-class mask model on PWMFD converted from VOC to YOLO format: 9,205 images, 18,528 labeled instances, 0.970 mAP@50 validation.

Traffic Violation Detection & License Plate Recognition

2026

Personal Project / Python, OpenCV, Ultralytics YOLO, ByteTrack, HyperLPR3, JSON / Source: [GitHub](#)

- Built an AI video analytics pipeline for vehicle detection/tracking, traffic-light state estimation, stop-line event logic, plate OCR, and JSON evidence bundles.
- Prepared a compact CCPD2019 plate dataset with annotation parsing, YOLO bbox conversion, train/val/test split, negative samples, OCR crops, and QA reports.
- Fine-tuned a YOLO plate detector on 8,598 images, achieving 0.989 precision, 0.998 recall, and 0.994 mAP50 on validation.

Autofall Labeler Tool

2026

Additional Project / Python, Streamlit, OpenCV, YOLO-Pose, Pandas / Source: [GitHub](#)

- Built a human-in-the-loop dataset assistant with Streamlit review UI, metadata/annotation management, YOLO-Pose keypoints, QA reports, and reviewed-only export.
- Designed CSV-based source-of-truth records for annotations, metadata, keypoints, and review events to handle noisy CCTV data and reproducible dataset versioning.
- Achieved 90.2% coarse fall/non-fall agreement on a reviewed benchmark while identifying exact action-label accuracy as a limitation of weak temporal modeling.

ACTIVITIES

AI Research Group – University Club

2025 – Present

Member

Discussed Machine Learning, Deep Learning, Computer Vision, and AI application development.